

```

In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
import numpy as np
import plotly.express as px
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_auc_score

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import RandomizedSearchCV
from xgboost import XGBClassifier as xgb

import shap
import lime
import lime.lime_tabular

```

```

In [2]: df = pd.read_csv("Iris Flower - Iris.csv")
df.head()

```

```

Out[2]:

```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [3]: df = df.drop('Id',axis=1)
df.head()
```

```
Out[3]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [4]: df1 =df.copy()
df2 =df.copy()
```

```
In [5]: df.shape
```

```
Out[5]: (150, 5)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   SepalLengthCm   150 non-null   float64
 1   SepalWidthCm    150 non-null   float64
 2   PetalLengthCm   150 non-null   float64
 3   PetalWidthCm    150 non-null   float64
 4   Species         150 non-null   object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

The target Variable is PetalLengthCm

```
In [8]: df.rename(columns={'Species': 'Species_Value'}, inplace=True)  
df.head()
```

```
Out[8]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

Chack Missing Values

```
In [9]: df.isnull().sum()
```

```
Out[9]: SepalLengthCm    0  
        SepalWidthCm     0  
        PetalLengthCm    0  
        PetalWidthCm     0  
        Species_Value    0  
        dtype: int64
```

```
In [10]: df.dtypes
```

```
Out[10]: SepalLengthCm    float64  
        SepalWidthCm     float64  
        PetalLengthCm    float64  
        PetalWidthCm     float64  
        Species_Value     object  
        dtype: object
```

Data inconsistencies & Data Preprocessing:

```
In [11]: for col in df.columns:
          print(f"---##### {col} ---#####")
          print(df[col].value_counts())
```

```
7.9      1
Name: SepalLengthCm, dtype: int64
---##### SepalWidthCm ---#####
3.0      26
2.8      14
3.2      13
3.1      12
3.4      12
2.9      10
2.7       9
2.5       8
3.5       6
3.3       6
3.8       6
2.6       5
2.3       4
3.7       3
2.4       3
2.2       3
~ ~ ~
```

Handling NaN Values

```
In [12]: Num_cols = []
          cat_cols = []
          for col in df.columns:
              if df[col].dtypes == 'object':
                  cat_cols.append(col)
              else:
                  Num_cols.append(col)
```

```
In [13]: #handle null values of numerical columns
```

```
for col in Num_cols:  
    if df[col].isna:  
        df[col].fillna(df[col].median(), inplace=True)
```

```
In [14]: #handle null values of categorical columns
```


```
for col in cat_cols:  
    if df[col].isnull:  
        df[col].fillna(df[col].mode()[0], inplace=True)
```

```
In [15]: df.isnull().sum()
```

```
Out[15]: SepalLengthCm    0  
SepalWidthCm             0  
PetalLengthCm            0  
PetalWidthCm             0  
Species_Value            0  
dtype: int64
```

Outliers

```
In [16]: plt.figure(figsize=(20,20))
for ax, col in enumerate(Num_cols):
    plt.subplot(5,6, int(ax+1))
    plt.title(col)
    sns.boxenplot(x=df[col],hue=df['Species_Value'])
plt.tight_layout()
plt.show()
```



Encoding target value

```
In [17]: from sklearn.preprocessing import LabelEncoder
led =LabelEncoder()
led.fit_transform(df['Species_Value'])
df['Species_Value'] =led.fit_transform(df['Species_Value'])
df.head()
```

```
Out[17]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [18]: df.value_counts('Species_Value')
```

```
Out[18]: Species_Value
0      50
1      50
2      50
dtype: int64
```

Model_Selection of ML

```
In [19]: df.head()
```

```
Out[19]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

Decision Tree Classifier

```
In [20]: from sklearn import tree
```

```
In [21]: df.head()
```

```
Out[21]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [22]: x =df.drop('Species_Value',axis=1)
y =df['Species_Value']
```

DecisionTreeClassifier

```
In [23]: from sklearn.tree import DecisionTreeClassifier
```

```
In [24]: dtc =DecisionTreeClassifier()
```

```
In [25]: dtc.fit(x,y)
```

```
Out[25]: DecisionTreeClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [26]: dtc.predict([[1,0,0,0]])
```

X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names

```
Out[26]: array([0])
```

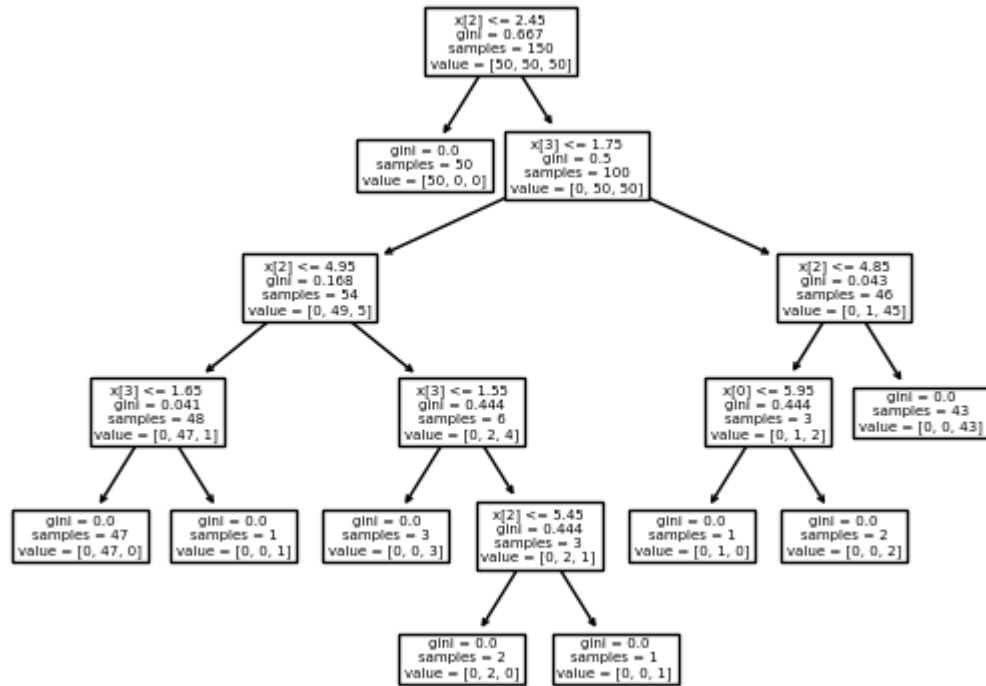
```
In [27]: dtc.predict([[0,0,0,1]])
```

X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names

```
Out[27]: array([0])
```

```
In [28]: tree.plot_tree(dtc)
```

```
Out[28]: [Text(0.5, 0.9166666666666666, 'x[2] <= 2.45\ngini = 0.667\nsamples = 150\nvalue = [50, 50, 50]'),
Text(0.4230769230769231, 0.75, 'gini = 0.0\nsamples = 50\nvalue = [50, 0, 0]'),
Text(0.5769230769230769, 0.75, 'x[3] <= 1.75\ngini = 0.5\nsamples = 100\nvalue = [0, 50, 50]'),
Text(0.3076923076923077, 0.5833333333333334, 'x[2] <= 4.95\ngini = 0.168\nsamples = 54\nvalue = [0, 4
9, 5]'),
Text(0.15384615384615385, 0.4166666666666667, 'x[3] <= 1.65\ngini = 0.041\nsamples = 48\nvalue = [0, 4
7, 1]'),
Text(0.07692307692307693, 0.25, 'gini = 0.0\nsamples = 47\nvalue = [0, 47, 0]'),
Text(0.23076923076923078, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.46153846153846156, 0.4166666666666667, 'x[3] <= 1.55\ngini = 0.444\nsamples = 6\nvalue = [0, 2,
4]'),
Text(0.38461538461538464, 0.25, 'gini = 0.0\nsamples = 3\nvalue = [0, 0, 3]'),
Text(0.5384615384615384, 0.25, 'x[2] <= 5.45\ngini = 0.444\nsamples = 3\nvalue = [0, 2, 1]'),
Text(0.46153846153846156, 0.08333333333333333, 'gini = 0.0\nsamples = 2\nvalue = [0, 2, 0]'),
Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.8461538461538461, 0.5833333333333334, 'x[2] <= 4.85\ngini = 0.043\nsamples = 46\nvalue = [0, 1,
45]'),
Text(0.7692307692307693, 0.4166666666666667, 'x[0] <= 5.95\ngini = 0.444\nsamples = 3\nvalue = [0, 1,
2]'),
Text(0.6923076923076923, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 1, 0]'),
Text(0.8461538461538461, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0, 0, 2]'),
Text(0.9230769230769231, 0.4166666666666667, 'gini = 0.0\nsamples = 43\nvalue = [0, 0, 43]')]
```



```
In [29]: tree.plot_tree(dtc, rounded=True, filled=True, feature_names = x.columns )
```

```
Out[29]: [Text(0.5, 0.9166666666666666, 'PetalLengthCm <= 2.45\ngini = 0.667\nsamples = 150\nvalue = [50, 50, 50]'),
Text(0.4230769230769231, 0.75, 'gini = 0.0\nsamples = 50\nvalue = [50, 0, 0]'),
Text(0.5769230769230769, 0.75, 'PetalWidthCm <= 1.75\ngini = 0.5\nsamples = 100\nvalue = [0, 50, 50]'),
Text(0.3076923076923077, 0.5833333333333334, 'PetalLengthCm <= 4.95\ngini = 0.168\nsamples = 54\nvalue = [0, 49, 5]'),
Text(0.15384615384615385, 0.4166666666666667, 'PetalWidthCm <= 1.65\ngini = 0.041\nsamples = 48\nvalue = [0, 47, 1]'),
Text(0.07692307692307693, 0.25, 'gini = 0.0\nsamples = 47\nvalue = [0, 47, 0]'),
Text(0.23076923076923078, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.46153846153846156, 0.4166666666666667, 'PetalWidthCm <= 1.55\ngini = 0.444\nsamples = 6\nvalue = [0, 2, 4]'),
Text(0.38461538461538464, 0.25, 'gini = 0.0\nsamples = 3\nvalue = [0, 0, 3]'),
Text(0.5384615384615384, 0.25, 'PetalLengthCm <= 5.45\ngini = 0.444\nsamples = 3\nvalue = [0, 2, 1]'),
Text(0.46153846153846156, 0.08333333333333333, 'gini = 0.0\nsamples = 2\nvalue = [0, 2, 0]'),
Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.8461538461538461, 0.5833333333333334, 'PetalLengthCm <= 4.85\ngini = 0.043\nsamples = 46\nvalue = [0, 1, 45]'),
Text(0.7692307692307693, 0.4166666666666667, 'SepalLengthCm <= 5.95\ngini = 0.444\nsamples = 3\nvalue = [0, 1, 2]'),
Text(0.6923076923076923, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 1, 0]'),
Text(0.8461538461538461, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0, 0, 2]'),
Text(0.9230769230769231, 0.4166666666666667, 'gini = 0.0\nsamples = 43\nvalue = [0, 0, 43]')]
```



K-Nearest Neighbors (KNN)

```
In [30]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [31]: knn = KNeighborsRegressor(n_neighbors=3)
```

```
In [32]: knn.fit(df[["Petal.Length.Cm"]],df["Species.Value"])
```

```
Out[32]: KNeighborsRegressor(n_neighbors=3)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


```
In [33]: knn.predict([[15]])
```

X does not have valid feature names, but KNeighborsRegressor was fitted with feature names

```
Out[33]: array([2.])
```

```
In [34]: plt.scatter(df.PetalWidthCm, df.Species_Value)  
plt.scatter(15,2, color="red")
```

```
Out[34]: <matplotlib.collections.PathCollection at 0x24d4af36510>
```



Logistic Regression

```
In [35]: df2.head()
```

```
Out[35]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [36]: df2.value_counts('Species')
```

```
Out[36]: Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
dtype: int64
```

```
In [37]: df2['Species'] = df2['Species'].replace(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'],[0,1,2])
```

```
In [38]: df2.head()
```

```
Out[38]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [39]: x=df2.drop('Species', axis=1)
```



```
In [47]: reg.predict_proba(x)
```

```
Out[47]: array([[9.81804276e-01, 1.81957100e-02, 1.43460484e-08],
 [9.71802160e-01, 2.81978107e-02, 2.97679380e-08],
 [9.85495978e-01, 1.45040094e-02, 1.21748649e-08],
 [9.76424964e-01, 2.35749969e-02, 3.91821962e-08],
 [9.85398356e-01, 1.46016317e-02, 1.18853668e-08],
 [9.70381261e-01, 2.96186651e-02, 7.36250444e-08],
 [9.86930605e-01, 1.30693757e-02, 1.97677029e-08],
 [9.76445341e-01, 2.35546320e-02, 2.74313737e-08],
 [9.79977769e-01, 2.00222013e-02, 3.01443225e-08],
 [9.69264678e-01, 3.07352902e-02, 3.12962705e-08],
 [9.76465697e-01, 2.35342841e-02, 1.92046470e-08],
 [9.75509897e-01, 2.44900592e-02, 4.34927557e-08],
 [9.74675031e-01, 2.53249478e-02, 2.11985210e-08],
 [9.92023518e-01, 7.97647851e-03, 3.82662375e-09],
 [9.88115480e-01, 1.18845174e-02, 2.82138058e-09],
 [9.86674116e-01, 1.33258709e-02, 1.29085668e-08],
 [9.88048548e-01, 1.19514428e-02, 9.20169859e-09],
 [9.81540546e-01, 1.84594348e-02, 1.95775987e-08],
 [9.56425841e-01, 4.35740903e-02, 6.85582969e-08],
 [9.84109627e-01, 1.58903526e-02, 2.04864764e-08],
 [9.46788727e-01, 5.32111874e-02, 8.60349466e-08],
 [9.81718341e-01, 1.82816258e-02, 3.27551617e-08],
 [9.96012118e-01, 3.98788090e-03, 1.30120720e-09],
 [9.52307406e-01, 4.76923584e-02, 2.35294370e-07],
 [9.52142739e-01, 4.78570566e-02, 2.04706858e-07],
 [9.51724920e-01, 4.82749940e-02, 8.58539787e-08],
 [9.69642321e-01, 3.03575937e-02, 8.57525714e-08],
 [9.74938912e-01, 2.50610632e-02, 2.48321606e-08],
 [9.77345878e-01, 2.26541046e-02, 1.73006443e-08],
 [9.71377478e-01, 2.86224638e-02, 5.78685246e-08],
 [9.64456751e-01, 3.55431792e-02, 6.96055318e-08],
 [9.64847917e-01, 3.51520261e-02, 5.72363014e-08],
 [9.88381497e-01, 1.16184958e-02, 7.03827774e-09],
 [9.89023113e-01, 1.09768821e-02, 5.32307431e-09],
 [9.69264678e-01, 3.07352902e-02, 3.12962705e-08],
 [9.84675735e-01, 1.53242569e-02, 7.90857558e-09],
 [9.78888779e-01, 2.11112114e-02, 9.59440027e-09],
 [9.69264678e-01, 3.07352902e-02, 3.12962705e-08],
 [9.85936993e-01, 1.40629917e-02, 1.52884991e-08],
 [9.74139741e-01, 2.58602304e-02, 2.82359399e-08],
 [9.86622239e-01, 1.33777494e-02, 1.12894387e-08],
 [9.62454256e-01, 3.75456797e-02, 6.46692339e-08],
 [9.89077657e-01, 1.09223321e-02, 1.11248863e-08],
```

[9.72415189e-01, 2.75846748e-02, 1.36473924e-07],
[9.60295084e-01, 3.97046941e-02, 2.22313454e-07],
[9.73941209e-01, 2.60587512e-02, 3.94699041e-08],
[9.80339423e-01, 1.96605521e-02, 2.52583405e-08],
[9.83423725e-01, 1.65762546e-02, 1.98937643e-08],
[9.78568486e-01, 2.14314958e-02, 1.86534512e-08],
[9.78710783e-01, 2.12891974e-02, 1.91086885e-08],
[2.10862668e-03, 8.73911720e-01, 1.23979654e-01],
[5.76823060e-03, 8.59784431e-01, 1.34447338e-01],
[1.05081764e-03, 7.25004743e-01, 2.73944440e-01],
[1.54491926e-02, 9.39605856e-01, 4.49449509e-02],
[2.36330137e-03, 8.15138141e-01, 1.82498557e-01],
[6.96159722e-03, 8.60065758e-01, 1.32972645e-01],
[3.73034137e-03, 7.16634303e-01, 2.79635356e-01],
[1.48642786e-01, 8.48302043e-01, 3.05517120e-03],
[2.76433824e-03, 8.96565673e-01, 1.00669988e-01],
[4.14097366e-02, 9.11964289e-01, 4.66259747e-02],
[5.63081930e-02, 9.37213785e-01, 6.47802175e-03],
[1.50880019e-02, 8.98923592e-01, 8.59884062e-02],
[9.14621362e-03, 9.76491713e-01, 1.43620734e-02],
[3.03185326e-03, 7.79195995e-01, 2.17772151e-01],
[7.43534625e-02, 9.15191810e-01, 1.04547279e-02],
[5.24771741e-03, 9.26352106e-01, 6.84001765e-02],
[8.65632963e-03, 7.74751407e-01, 2.16592263e-01],
[1.64699180e-02, 9.65138181e-01, 1.83919005e-02],
[1.80985911e-03, 8.00836957e-01, 1.97353184e-01],
[2.40284995e-02, 9.59352815e-01, 1.66186855e-02],
[2.27943755e-03, 4.40397579e-01, 5.57322984e-01],
[1.67935233e-02, 9.56795838e-01, 2.64106382e-02],
[7.12108592e-04, 5.96065484e-01, 4.03222407e-01],
[3.03114082e-03, 8.59773974e-01, 1.37194885e-01],
[7.04928838e-03, 9.42941928e-01, 5.00087837e-02],
[5.04702753e-03, 9.20090125e-01, 7.48628470e-02],
[1.11365773e-03, 8.01376342e-01, 1.97510000e-01],
[5.72573050e-04, 4.81147340e-01, 5.18280087e-01],
[5.45016575e-03, 8.13094564e-01, 1.81455270e-01],
[6.19098607e-02, 9.34706775e-01, 3.38336387e-03],
[2.92562226e-02, 9.57114156e-01, 1.36296213e-02],
[3.73398416e-02, 9.55117128e-01, 7.54303014e-03],
[2.51785108e-02, 9.56466417e-01, 1.83550721e-02],
[4.46077656e-04, 3.49641519e-01, 6.49912404e-01],
[1.01618347e-02, 7.50991852e-01, 2.38846313e-01],
[9.88064442e-03, 7.88903270e-01, 2.01216085e-01],

[2.24498796e-03, 8.05140196e-01, 1.92614816e-01],
[2.76574182e-03, 9.13049949e-01, 8.41843090e-02],
[2.69596896e-02, 9.28603489e-01, 4.44368212e-02],
[1.99216288e-02, 9.38039653e-01, 4.20387180e-02],
[8.71455579e-03, 8.97810703e-01, 9.34747407e-02],
[4.61617435e-03, 8.28235221e-01, 1.67148604e-01],
[1.75871757e-02, 9.56979931e-01, 2.54328934e-02],
[1.22509266e-01, 8.74440260e-01, 3.05047361e-03],
[1.44435519e-02, 9.20449668e-01, 6.51067802e-02],
[1.99267378e-02, 9.38131157e-01, 4.19421055e-02],
[1.70447215e-02, 9.25454598e-01, 5.75006806e-02],
[8.46623899e-03, 9.35114317e-01, 5.64194440e-02],
[2.44641374e-01, 7.54068127e-01, 1.29049945e-03],
[1.91187811e-02, 9.36056233e-01, 4.48249863e-02],
[8.83945976e-07, 3.92380755e-03, 9.96075309e-01],
[2.40478186e-04, 1.62637592e-01, 8.37121929e-01],
[2.43989458e-06, 2.55900308e-02, 9.74407529e-01],
[3.08462316e-05, 8.17010657e-02, 9.18268088e-01],
[3.67341536e-06, 1.74640166e-02, 9.82532310e-01],
[5.42581633e-08, 4.64029239e-03, 9.95359653e-01],
[5.74744241e-03, 5.13824824e-01, 4.80427733e-01],
[6.12087217e-07, 2.13369896e-02, 9.78662398e-01],
[5.17296248e-06, 5.32566160e-02, 9.46738211e-01],
[6.38000895e-07, 5.75609365e-03, 9.94243268e-01],
[2.97721865e-04, 2.10640717e-01, 7.89061562e-01],
[7.19637222e-05, 1.37323809e-01, 8.62604227e-01],
[2.09423186e-05, 6.53123413e-02, 9.34666716e-01],
[2.27360977e-04, 1.45309102e-01, 8.54463537e-01],
[6.80627890e-05, 4.35826418e-02, 9.56349295e-01],
[5.07181450e-05, 5.41253766e-02, 9.45823905e-01],
[5.48969877e-05, 1.22934156e-01, 8.77010947e-01],
[8.24635299e-08, 3.56625143e-03, 9.96433666e-01],
[3.09054643e-09, 1.00046829e-03, 9.98999529e-01],
[3.88669705e-04, 4.51638362e-01, 5.47972968e-01],
[5.49480546e-06, 2.38827666e-02, 9.76111739e-01],
[6.07239782e-04, 1.90609152e-01, 8.08783608e-01],
[3.08193167e-08, 4.65327551e-03, 9.95346694e-01],
[5.80434319e-04, 3.93056939e-01, 6.06362627e-01],
[1.26038590e-05, 3.86515075e-02, 9.61335889e-01],
[4.77771806e-06, 5.14939827e-02, 9.48501240e-01],
[1.06269412e-03, 4.56521122e-01, 5.42416184e-01],
[1.01199819e-03, 3.85479491e-01, 6.13508511e-01],
[1.05051565e-05, 3.63647722e-02, 9.63624723e-01],

```
[1.66741235e-05, 1.41905719e-01, 8.58077606e-01],  
[1.04799974e-06, 2.91884542e-02, 9.70810498e-01],  
[6.86347607e-07, 1.74357944e-02, 9.82563519e-01],  
[7.76839814e-06, 2.72882629e-02, 9.72703969e-01],  
[5.24875293e-04, 4.75421735e-01, 5.24053390e-01],  
[6.22412534e-05, 1.88447131e-01, 8.11490627e-01],  
[3.85676613e-07, 1.17528221e-02, 9.88246792e-01],  
[1.13862839e-05, 1.73678116e-02, 9.82620802e-01],  
[6.67619232e-05, 1.19535901e-01, 8.80397337e-01],  
[1.60407457e-03, 4.40522497e-01, 5.57873429e-01],  
[3.90465598e-05, 9.35554232e-02, 9.06405530e-01],  
[6.18835283e-06, 2.03228444e-02, 9.79670967e-01],  
[9.80090824e-05, 1.20781073e-01, 8.79120918e-01],  
[2.40478186e-04, 1.62637592e-01, 8.37121929e-01],  
[2.00921845e-06, 1.26057079e-02, 9.87392283e-01],  
[3.73036310e-06, 1.21304884e-02, 9.87865781e-01],  
[5.50264541e-05, 8.02057758e-02, 9.19739198e-01],  
[2.25169948e-04, 2.51893935e-01, 7.47880895e-01],  
[1.36503465e-04, 1.57222283e-01, 8.42641214e-01],  
[4.47860511e-05, 3.85093811e-02, 9.61445833e-01],  
[4.70348664e-04, 2.34999252e-01, 7.64530399e-01]])
```

```
In [48]: xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=.3, random_state=42)
```

LogisticRegression

```
In [49]: LR_model = LogisticRegression()
```

```
In [50]: LR_model.fit(xtrain, ytrain)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
Out[50]: LogisticRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [51]: smote = SMOTE(random_state=42)
xtrain, ytrain = smote.fit_resample(xtrain, ytrain)
```

RandomForestClassifier

```
In [52]: RFC = RandomForestClassifier()
```

```
In [53]: RFC.fit(xtrain, ytrain)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
Out[53]: RandomForestClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

XGboost

```
In [54]: XGB = xgb()
```

```
In [55]: XGB.fit(xtrain, ytrain)
```

```
Out[55]: XGBClassifier(base_score=None, booster=None, callbacks=None,  
    colsample_bylevel=None, colsample_bynode=None,  
    colsample_bytree=None, device=None, early_stopping_rounds=None,  
    enable_categorical=False, eval_metric=None, feature_types=None,  
    gamma=None, grow_policy=None, importance_type=None,  
    interaction_constraints=None, learning_rate=None, max_bin=None,  
    max_cat_threshold=None, max_cat_to_onehot=None,  
    max_delta_step=None, max_depth=None, max_leaves=None,  
    min_child_weight=None, missing=nan, monotone_constraints=None,  
    multi_strategy=None, n_estimators=None, n_jobs=None,  
    num_parallel_tree=None, objective='multi:softprob', ...)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

SVC

```
In [56]: svc = SVC(kernel='rbf', gamma='scale', probability=True)
```

```
In [57]: svc.fit(xtrain, ytrain)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
Out[57]: SVC(probability=True)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

Predicting and Evaluation Metrics


```
In [58]: trained_models = [  
    LR_model,  
    RFC,  
    XGB,  
    svc  
]
```

```

In [59]: for model in trained_models:
    y_pred = model.predict(xtest)
    accuracy = accuracy_score(ytest, y_pred)
    precision = precision_score(ytest, y_pred, average='micro')
    recall = recall_score(ytest, y_pred, average='micro')
    f1 = f1_score(ytest, y_pred, average='micro')
    cm = confusion_matrix(ytest, y_pred)

    proba = model.predict_proba(xtest)
    auc = roc_auc_score(ytest, proba, multi_class='ovr')

    print(f"\t ***** {model.__class__.__name__} *****")
    print("Accuracy Score: ", accuracy)
    print("Precision Score: ", precision)
    print("Recall Score: ", recall)
    print("F1 Score: ", f1)
    print("AUC Score: ", auc)
    print("Confusion Matrix: \n", cm)
    sns.heatmap(cm)
    plt.show()

    fpr, tpr, threshold = roc_curve(ytest, proba[:,1], pos_label=1)
    plt.plot(fpr, tpr, linestyle="--", label="CURVE", )
    plt.title("ROC CURVE")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.legend()
    plt.show()

```

***** LogisticRegression *****

```

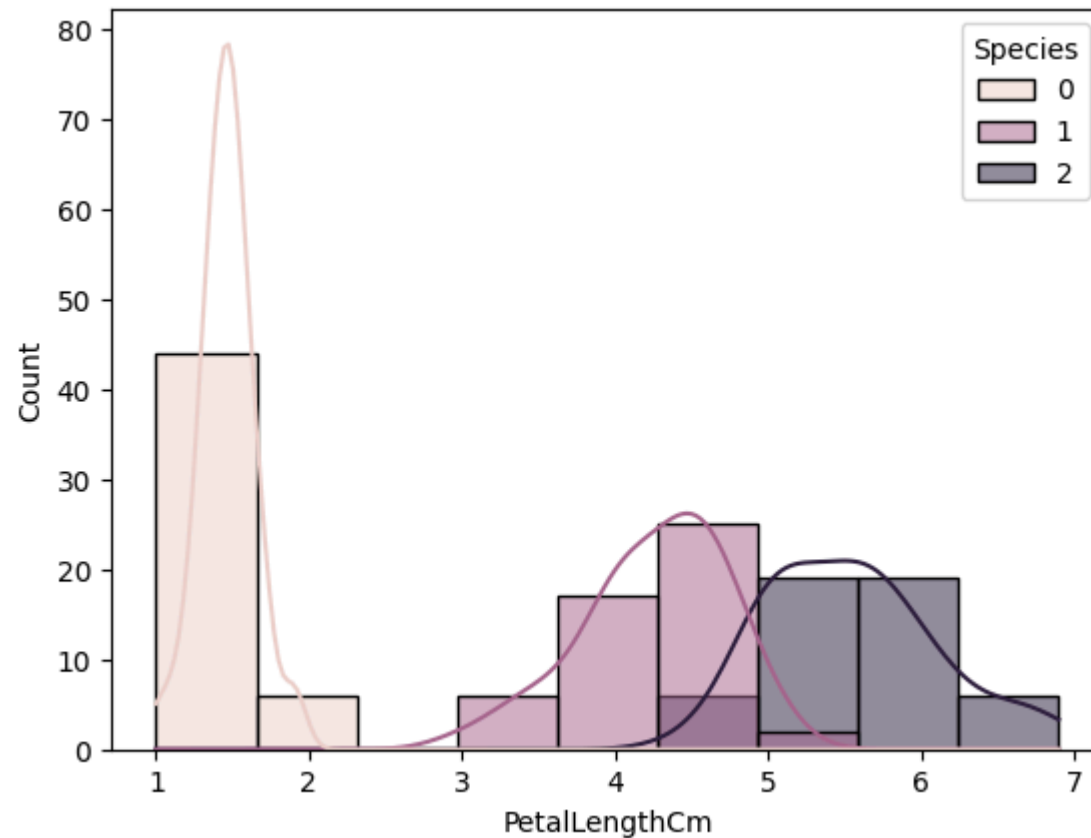
Accuracy Score:  1.0
Precision Score:  1.0
Recall Score:    1.0
F1 Score:        1.0
AUC Score:       1.0
Confusion Matrix:
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]

```



```
In [60]: sns.histplot(x = "PetalLengthCm", data = df2, kde=True, hue="Species")
```

```
Out[60]: <Axes: xlabel='PetalLengthCm', ylabel='Count'>
```



RandomForest Classifier has the highest accuracy. So we will use Gradient Boosting Classifier for hyper parameter tuning

Hyperparameter Tuning:

```
In [61]: params = {  
    'n_estimators': [100, 200, 300],  
    'criterion': ['gini', 'entropy'],  
    'max_depth': [None, 10, 20, 30],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 3],  
    'max_features': ['auto', 'sqrt', 'log2']  
}
```

```
In [62]: model_param = {
    'svm': {
        'model': SVC(gamma='auto'),
        'params': {
            'C': [1.0, 5.0, 10.0],
            'kernel': ['rbf', 'linear']
        }
    },
    'LogReg': {
        'model': LogisticRegression(solver='liblinear'),
        'params': {
            'C': [1.0, 5.0, 10.0],
            'penalty': ['l1', 'l2'],
        }
    },
    'rf': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [20, 50, 100],
            'criterion': ['gini', 'entropy'],
            'min_samples_leaf': [1, 2],
            'max_features': ['sqrt', 'log2']
        }
    },
    'XGboost': {
        'model': xgb(),
        'params': {
            'learning_rate': [0.1, 0.01, 0.2],
            'n_estimators': [20, 50, 100],
        }
    }
}
```

```
In [63]: scores = []

for name, mp in model_param.items():
    RandomSearch = RandomizedSearchCV(estimator=mp['model'] , param_distributions=mp['params'], return_tr
    RandomSearch.fit(xtrain, ytrain)
    scores.append({
        'model': mp['model'],
        'best_score': RandomSearch.best_score_,
        'best_param': RandomSearch.best_params_
    })
```

[illegible]

```
In [64]: scores
```

```
Out[64]: [{'model': SVC(gamma='auto'),
  'best_score': 0.9549407114624506,
  'best_param': {'kernel': 'rbf', 'C': 10.0}},
 {'model': LogisticRegression(solver='liblinear'),
  'best_score': 0.9549407114624506,
  'best_param': {'penalty': 'l2', 'C': 10.0}},
 {'model': RandomForestClassifier(),
  'best_score': 0.9367588932806324,
  'best_param': {'n_estimators': 50,
  'min_samples_leaf': 1,
  'max_features': 'sqrt',
  'criterion': 'entropy'}},
 {'model': XGBClassifier(base_score=None, booster=None, callbacks=None,
  colsample_bylevel=None, colsample_bynode=None,
  colsample_bytree=None, device=None, early_stopping_rounds=None,
  enable_categorical=False, eval_metric=None, feature_types=None,
  gamma=None, grow_policy=None, importance_type=None,
  interaction_constraints=None, learning_rate=None, max_bin=None,
  max_cat_threshold=None, max_cat_to_onehot=None,
  max_delta_step=None, max_depth=None, max_leaves=None,
  min_child_weight=None, missing=nan, monotone_constraints=None,
  multi_strategy=None, n_estimators=None, n_jobs=None,
  num_parallel_tree=None, random_state=None, ...),
  'best_score': 0.9276679841897234,
  'best_param': {'n_estimators': 100, 'learning_rate': 0.2}}]
```

Predict testing values with hyperparameter tuning

```
In [65]: rf = RandomForestClassifier(n_estimators=50, min_samples_leaf=1, criterion='entropy')
rf.fit(xtrain, ytrain)
b_predicted = rf.predict(xtest)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_sample s,), for example using ravel().


```
In [66]: accuracy_score(ytest, b_predicted)
```

```
Out[66]: 1.0
```


```
In [67]: accuracy = accuracy_score(ytest, b_predicted)
precision = precision_score(ytest, b_predicted, average='micro')
recall = recall_score(ytest, b_predicted, average='micro')
f1 = f1_score(ytest, b_predicted, average='micro')
cm = confusion_matrix(ytest, b_predicted)
```

```
In [68]: print("Accuracy Score: ", accuracy)
print("Precision Score: ", precision)
print("Recall Score: ", recall)
print("F1 Score: ", f1)
print("AUC Score: ", auc)
print("Confusion Matrix: \n", cm)
```

```
Accuracy Score: 1.0
Precision Score: 1.0
Recall Score: 1.0
F1 Score: 1.0
AUC Score: 1.0
Confusion Matrix:
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```


```
In [69]: sns.heatmap(cm)
```

```
Out[69]: <Axes: >
```



Interpretability:

```
In [70]: explainer = shap.Explainer(rf)
shap_values = explainer.shap_values(xtest)
shap.summary_plot(shap_values, xtest)
```



```
In [ ]:
```